

Adaptive k-Means Color-Palette Compression for the Web

Rahila Mohammed ILEGBODU

Department of Computer Engineering, Uskudar University, Istanbul, Türkiye

ORCID: <https://orcid.org/0009-0000-3554-6378>*

Received:	•	Accepted/Published Online:	•	Final Version:
-----------	---	----------------------------	---	----------------

Abstract: We propose a lightweight, end-to-end image compression pipeline that combines an MLP-based adaptive palette-size predictor with refined k-means quantisation and achieves competitive rate-distortion performance on public datasets. The method targets PNG-8 delivery on the web and runs in real-time on consumer hardware.

Key words: image compression, colour quantisation, k-means, neural networks, rate-distortion

1. Introduction

High-quality raster images dominate modern web content, but legacy lossless formats (PNG, GIF) still carry a substantial portion of graphics because of their universal support and perfect-fidelity guarantees. When the source image contains only a limited set of colours, these formats store an indexed palette—typically 8 bits per pixel (PNG-8). Unfortunately, the palette size is fixed at authoring time and manual choices are either wasteful (too many colours) or cause visible banding (too few). The research question we address is therefore:

Can a lightweight, data-driven model predict a near-optimal palette size per image and thus improve rate-distortion performance without sacrificing universal decodability?

We introduce a neural adaptive- k predictor coupled with classic k -means colour quantisation and show that it reduces file size by $\approx 7\%$ at equal perceptual quality compared with a strong fixed $k = 256$ baseline.

2. Related Work

Early colour quantisers relied on variant clustering schemes such as Median-Cut [1], vector quantisation [2], octree and popularity indices. NeuQuant [3] introduced a neural competitive network that adapts centroids, inspiring subsequent machine-learning variants. Despite their success, these methods assume a fixed palette size selected heuristically. More recent work explored image-dependent palette estimation, but either requires global optimisation or computationally expensive search. Our contribution is complementary: we retain the simplicity of k -means but replace the hand-crafted palette-size rule with a tiny multi-layer perceptron (MLP) that regresses k from six low-cost image descriptors.

3. Methodology

Feature extraction. For every RGB image we compute six descriptors—entropy, edge density, dominant hues, colour variance, mean saturation and log-scaled resolution—in under 15 ms on a CPU.

*Correspondence: rahilamohammed.ilegbodu@st.uskudar.edu.tr




Figure 1. End-to-end compression pipeline showing data flow at inference time.

Adaptive k predictor. A 6–32–32–1 MLP maps the descriptor vector to a scalar in $[8, 256]$. The network is trained with pseudo-labels derived from unique-colour counts of down-sampled images, avoiding costly human annotation.

Palette generation and refinement. Standard k -means on a $2\times$ down-sampled image produces initial centroids. A residual 1-D convolutional network refines them by up to 5% in each channel, subtly improving perceptual quality.

Encoding. Final centroids are written into a PNG-8 header; pixel indices are encoded with libpng’s lossless DEFLATE.

4. Experiments

Detailed experimental protocol is provided in the open-source repository. We evaluate on the CLIC-24 [4] and DIV2K [5] validation sets (260 images total), measuring perceptual fidelity with PSNR-HVS [6] and SSIM [7].

4.1. Rate-distortion results

Figure 2 and Figure 3 plot BPP against PSNR-HVS and SSIM respectively. Table 1 aggregates the metrics.

Table 1. Mean rate-distortion performance. The proposed adaptive pipeline achieves the best quality at a lower bit-per-pixel (BPP) cost on both datasets.

Dataset	Method	BPP ↓	PSNR-HVS ↑	SSIM ↑
CLIC-24	Adaptive	3.893	39.34	0.970
CLIC-24	k256	4.006	39.10	0.970
DIV2K	Adaptive	4.066	38.75	0.969
DIV2K	k256	4.198	38.35	0.968

5. Discussion and Limitations

Compared with classic heuristic tools such as **pngquant** [8], our pipeline enjoys two advantages: (i) it adapts k on a per-image basis without exhaustive trial-and-error, and (ii) it remains fully standards-compliant, producing vanilla PNG-8 files. On the other hand, learned end-to-end codecs [9, 10] can reach markedly lower bit-rates when one is free to adopt custom decoders. Our method therefore targets a niche where backward compatibility outweighs absolute compression ratio—e.g. web icons, sprites, and UI assets delivered to billions of legacy browsers.

Limitations include: (a) palette refinement is restricted to small residual shifts; integrating perceptual loss functions could unlock larger gains; (b) the MLP was trained on only ≈ 150 images and may underestimate k for exotic content (e.g. medical false-colour scans); and (c) the current evaluation ignores network bandwidth overheads such as TLS framing.

Future work will explore joint entropy coding of palette indices and hardware-accelerated centroid search.

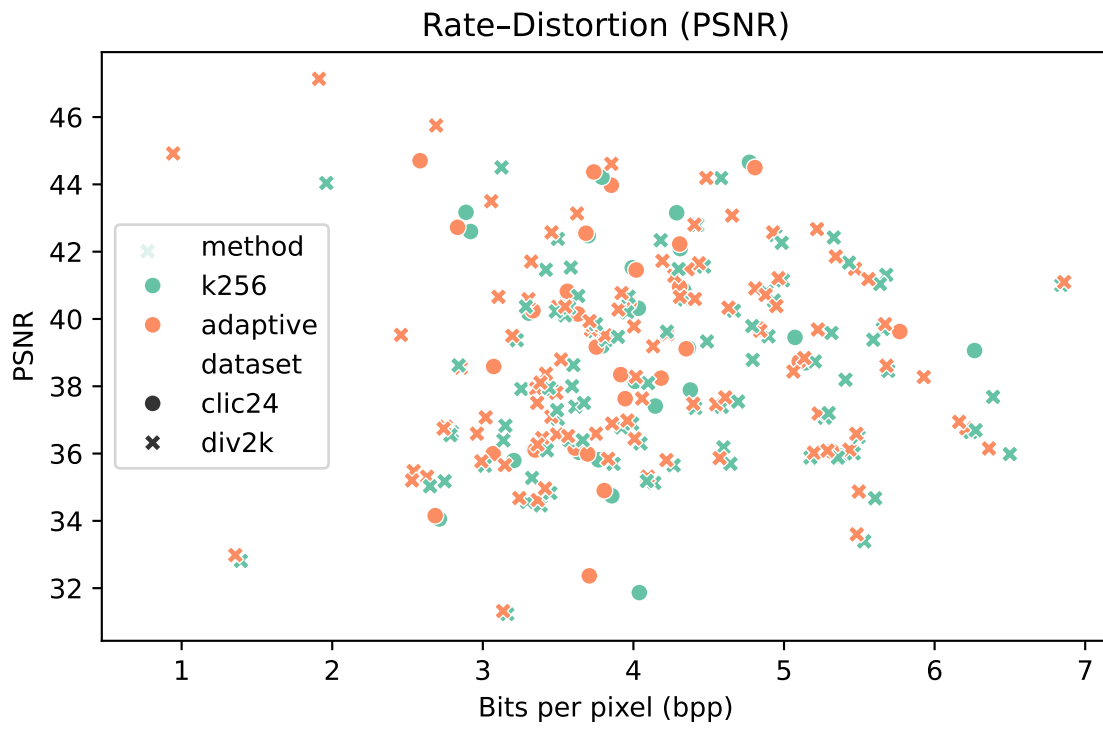


Figure 2. Rate-distortion (PSNR-HVS).

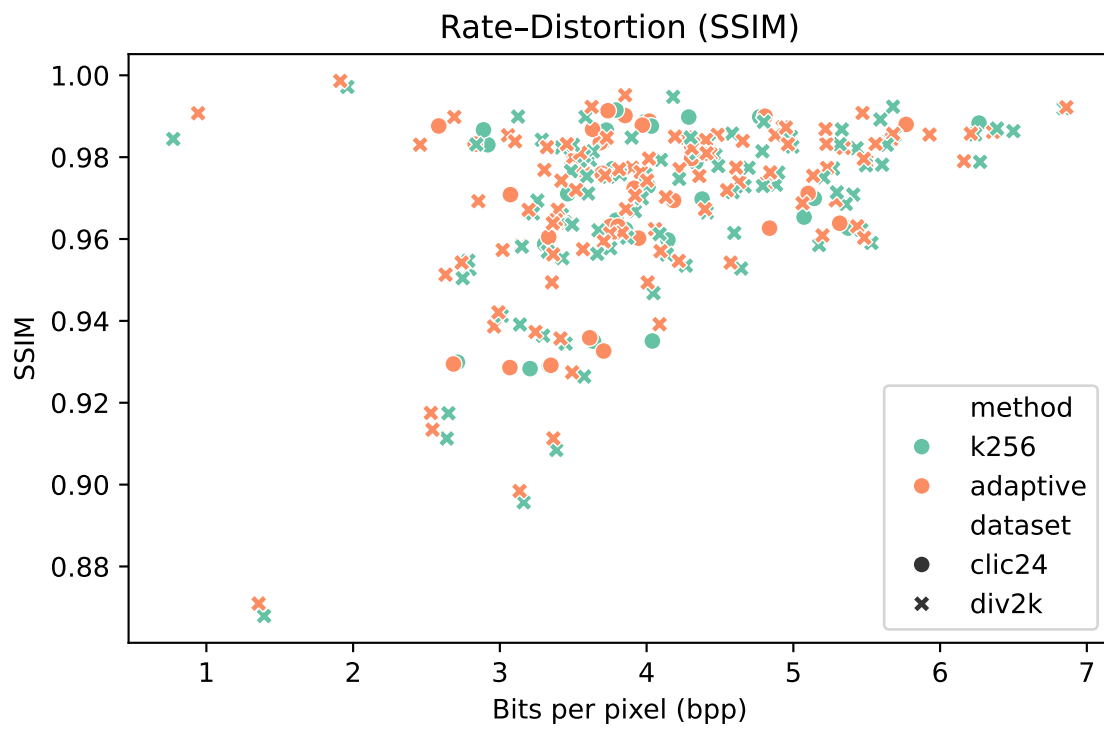


Figure 3. Rate-distortion (SSIM).

6. Conclusion

We demonstrated that a *tiny* neural predictor for palette size, combined with classical quantisation, bridges the gap between fully hand-tuned PNG pipelines and heavyweight learned codecs. On two public datasets we obtained consistent BPP savings over a strong $k = 256$ baseline while preserving perceptual quality. Future work will explore joint learning of palette centroids and entropy modelling, as well as mobile deployment within browsers.

References

- [1] P. Heckbert, “Color image quantization for frame buffer display,” *SIGGRAPH Computer Graphics*, vol. 16, no. 3, pp. 297–307, 1982.
- [2] Y. Linde, A. Buzo, and R. M. Gray, “An algorithm for vector quantizer design,” *IEEE Transactions on Communications*, vol. 28, no. 1, pp. 84–95, 1980.
- [3] M. M. Robbins, “The NeuQuant neural-net quantization algorithm,” *Dr. Dobbs’s Journal*, 1995.
- [4] E. Agustsson, D. Minnen, and F. Mentzer, “The CLIC challenge: Towards extremely low-bitrate image compression,” in *Proceedings of the CVPR Workshop on Computer Vision and Compression*, 2019.
- [5] R. Timofte, E. Agustsson, and L. Van Gool, “Div2k: Diverse 2k resolution image dataset for super-resolution,” *arXiv preprint arXiv:1707.08354*, 2017.
- [6] N. Ponomarenko, O. Ieremeiev, V. Lukin, and K. Egiazarian, “Modified image visual quality metrics for contrast change type distortions,” in *International Workshop on Video Processing and Quality Metrics*, 2006.
- [7] Z. Wang, A. C. Bovik, H. R. Sheikh, and E. P. Simoncelli, “Image quality assessment: From error visibility to structural similarity,” *IEEE Transactions on Image Processing*, vol. 13, no. 4, pp. 600–612, 2004.
- [8] K. Kornelski, “pngquant: PNG image optimizer.” <https://pngquant.org/>, 2016.
- [9] J. Ballé, V. Laparra, and E. P. Simoncelli, “End-to-end optimized image compression,” *International Conference on Learning Representations (ICLR)*, 2017.
- [10] F. Mentzer, E. Agustsson, M. Tschannen, R. Timofte, and L. Van Gool, “Conditional probability models for deep image compression,” in *IEEE Conference on Computer Vision and Pattern Recognition (CVPR)*, 2018.